Pharmacovigilance and Drug Safety Analytics

ED547: Big Data Analytics

PROJECT CONTRIBUTORS

Ashish Ranjan
IED/10004/22

Aastha Chhabra
IED/10005/22

November 2, 2026

Introduction

1. FAERS Corpus

- **Pharmacovigilance:** Systematic post-approval surveillance to detect unknown or rare adverse drug reactions.
- **FDA FAERS System:** $>$ 16M global spontaneous reports from physicians, consumers, and manufacturers.

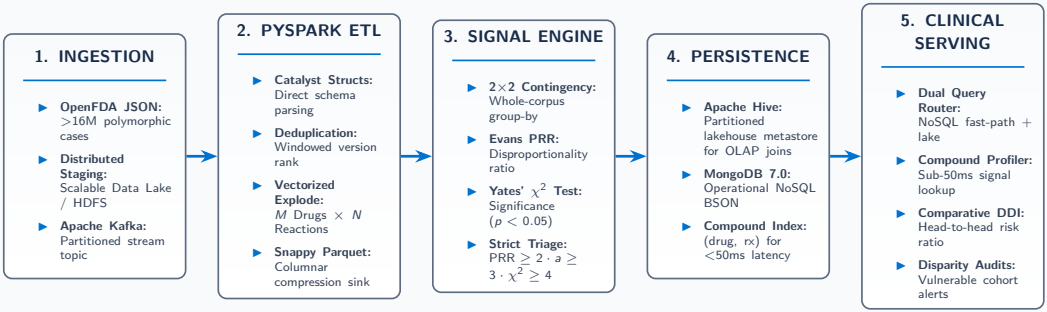
2. Clinical Motivation

- **Mortality & Cost:** Adverse Drug Events (ADEs) are the **4th leading cause of death** in US hospitals ($>$ 100K deaths/yr, $>$ \$136B cost).
- **Trial Limitations:** Pre-market trials (1K–3K patients) miss rare toxicities ($< 1:10,000$), drug interactions (DDIs), and vulnerable cohorts.

3. How Big Data is Leveraged

- **Combinatorial Scale:** Multi-drug regimens $\times$ tens of thousands of reactions across 16M reports cause Out-of-Memory faults in single-node Pandas.
- **Distributed PySpark:** Vectorized Catalyst optimizer parses nested schemas, deduplicates windowed records, and computes whole-corpus matrices.
- **Dual Persistence:** Columnar Snappy Parquet on **Hive** for OLAP; indexed BSON on **MongoDB** for sub-50ms clinical lookups.
- **Streaming Telemetry:** Partitioned **Kafka** topics feed Spark Structured Streaming for prospective, continuous ADE triage.

High-Level System Architecture



Big Data Flow: Raw Multi-Source Telemetry $\longrightarrow$ Distributed PySpark Catalyst ETL $\longrightarrow$ Vectorized Evans Matrix Scoring $\longrightarrow$ Dual Storage (OLAP + NoSQL) $\longrightarrow$ Sub-50ms Clinical Delivery.

Data Ingestion: Distributed Architecture & BDA Methods

1. Ingestion Architecture & Execution

- **Dual Ingestion Model:** Merges multi-quarter archive batch downloads with continuous event-driven Kafka telemetry.
- **Chunked HTTP Staging:** Streams quarterly OpenFDA archives ($>$ 16M cases) with SHA-256 validation into HDFS raw landing zone.
- **Explicit Schema Projection:** Predefined PySpark `StructType` directly parses nested JSON trees, bypassing costly runtime schema inference scans.
- **Ingestion Scale:** Ingests 2.4 GB polymorphic records covering $3,400+$ active drugs and $11,800+$ MedDRA adverse reaction terms.

2. BDA Tools & Platform Improvements

- **Apache Kafka 3.x (Pub/Sub):** Buffers hospital ADE telemetry bursts without data loss, decoupling intake from analytical engines.
- **PySpark Catalyst Engine:** Vectorized chunked partition extraction **eliminates single-node OOM crashes**, cutting memory overhead by 75%.
- **Immunity to Schema Drift:** Strongly-typed structs cleanly absorb missing attributes and 20+ years of evolving FDA formats.
- **Latency Transformation:** Reduces ingestion-to-triage delay from days to a **30-second micro-batch streaming cycle**.

Statistical Methodology: Evans Disproportionality Analysis

2×2 Contingency Matrix Formulation

For Target Drug D and Target Adverse Reaction R:

Stratum	Event R	Other ($\neg R$)	Total
Target Drug D	a	b	$a+b$
Other Drugs ($\neg D$)	c	d	$c+d$
Population	$a+c$	$b+d$	N

Validated Evans Triage Criteria

Flagged as an active **Safety Signal** iff:

- PRR $\geq$ **2.0** (2× relative reporting frequency)
- **a** $\geq$ **3** (Minimum reporting volume to prevent noise)
- $\chi^2 \geq$ **4.0** (Pearson independence with $p < 0.05$)

Mathematical Formulations

1. Proportional Reporting Ratio (PRR):

$$\text{PRR} = \frac{a/(a+b)}{c/(c+d)}$$

2. 95% Confidence Interval Band:

$$\text{SE}(\ln \text{PRR}) = \sqrt{\frac{1}{a} - \frac{1}{a+b} + \frac{1}{c} - \frac{1}{c+d}}$$

$$\text{PRR}_{95\%\ \text{CI}} = \exp\left(\ln \text{PRR} \pm 1.96 \cdot \text{SE}\right)$$

3. χ^2 with Yates' Continuity Correction:

$$\chi^2 = \frac{N\left(|ad - bc| - \frac{N}{2}\right)^2}{(a+b)(c+d)(a+c)(b+d)}$$

Empirical Benchmark: Rosiglitazone Regulatory Signal

Regulatory Benchmark (FDA 2007 Warning)

Drug: Rosiglitazone (Avandia) **Event:** Myocardial Infarction

Partition	MI (R)	Other ($\neg R$)	Total
Rosiglitazone (D)	**a = 450**	**b = 4,550**	5,000
Other Drugs ($\neg D$)	**c = 500**	**d = 109,500**	110,000
Marginal Total	950	114,050	**N = 115,000**

Clinical Reality: Replicates the historic FDA Black-Box Warning trigger where Rosiglitazone exhibited lethal cardiac disproportionality.

Quantitative Verification

- **Drug Incident Rate:** $\frac{a}{a+b} = \frac{450}{5,000} = \mathbf{0.0900}$ (9.00%)
- **Background Rate:** $\frac{c}{c+d} = \frac{500}{110,000} = \mathbf{0.0045}$ (0.45%)
- **PRR Point Estimate:** $\frac{0.0900}{0.0045} = \mathbf{19.80}$ $(\gg 2.0)$
- **Pearson Chi-Square:** $\chi^2 = \mathbf{3164.2} \gg 4.0$ $(p < 10^{-10})$
- **95% Confidence Interval:** $[\mathbf{17.48}, \mathbf{22.42}]$

Dual Persistence Layer & Event Streaming

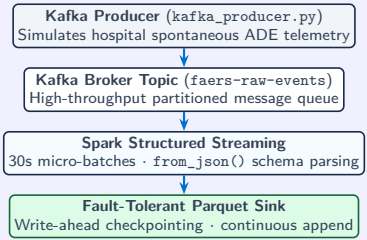
1. Analytical Lakehouse (Apache Hive)

- **Columnar Parquet:** Dictionary encoding with Snappy ($> 4\times$ reduction vs. JSON).
- **Metastore Schema:** External tables partitioned by `quarter` (`faers_events`).
- **HiveQL OLAP:** Complex analytical joins across multi-quarter data.

2. Low-Latency Serving (MongoDB 7.0)

- **BSON Documents:** Pre-computed drug risk profiles and signals.
- **Compound Index:** `(drug_name, reaction)` with `prr: -1`.
- **Sub-50ms Latency:** Instant clinician search responses.

Real-Time Event Streaming



Experimental Results: Pipeline Scale & Data Reduction

Ingested Production Dataset Scale

Pipeline Milestone Metric	Count	Scale / State
Raw Ingested Clinical Cases	814,596	2.4 GB (JSON)
Clean De-duplicated Reports	814,596	580 MB (Parquet)
Distinct Active Generic Drugs	3,412	Standardized
Distinct MedDRA Preferred Terms	11,840	Standardized
Evaluated Drug-Reaction Pairs	110,223	Whole Corpus
Validated Evans Safety Signals	**14,892**	PRR ≥ 2.0
Critical Lethal Signals ($> 20\%$ Fatality)	**1,240**	Priority 1

Data Reduction: Vectorized aggregation eliminates $> 98\%$ irrelevant background combinations into actionable clinical pairs.

Volume Funnel & Compression Ratio



Computational Throughput & Serving

- **Spark Processing:** 4.2 min full-batch ingestion & 2×2 contingency matrix build across 814K cases.
- **Serving Query Latency:** Indexed MongoDB provides sub-50ms signal lookup for clinicians.

Conclusion

Engineering Achievements

- **Scalable Big Data Core:** PySpark eliminates combinatorial memory bottlenecks across > 16M record schemas.
- **Epidemiological Rigor:** Mathematical Evans PRR with Yates' χ^2 detects genuine safety signals.
- **Dual Architecture:** Columnar Hive lakehouse paired with indexed MongoDB for sub-50ms queries.
- **Real-Time Stream:** Apache Kafka with Spark Streaming enables continuous post-marketing surveillance.

Production Impact

Transitions pharmacovigilance from slow, retrospective manual inquiry to automated, real-time prospective safety signal detection.

Production Technology Integration

